

```

In [1]: ▶ 1 #Importing Libraries needed
          2 import pandas as pd
          3 import numpy as np
          4
          5
          6 #For visualization
          7 import matplotlib.pyplot as plt
          8 from matplotlib import pyplot
          9 %matplotlib inline
         10 import seaborn as sns
         11 import shap
         12
         13 # For our modeling steps
         14 from sklearn.model_selection import train_test_split
         15 from sklearn.preprocessing import normalize
         16 from sklearn.linear_model import LinearRegression, LogisticRegression
         17 from sklearn.metrics import log_loss
         18 from sklearn.metrics import precision_recall_fscore_support
         19 from sklearn.preprocessing import StandardScaler
         20 from sklearn.svm import SVC
         21 from sklearn.metrics import confusion_matrix
         22 from sklearn.metrics import precision_score, recall_score, f1_score, a
         23 from sklearn.model_selection import GridSearchCV, cross_val_score
         24 import xgboost as xgb
         25 from xgboost import XGBClassifier, plot_importance
         26 from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_
         27 from sklearn import svm, datasets
         28 from sklearn import metrics
         29
         30
         31 # For demonstrative purposes
         32 from scipy.special import logit, expit
         33
         34 import warnings
         35 warnings.filterwarnings('ignore')

```

pandas.Int64Index is deprecated and will be removed from pandas in a future version. Use pandas.Index with the appropriate dtype instead.

How Do Wind Speed Features Impact Model Performance

For the scope of this project we are particularly interested in how hurricane's impact real estate value. Using wind speed we were able to gauge how much damage a hurricane had on a certain area. The most important feature for our best model was the fastest 2 minute wind gust and SizeRank. When we drop the wind features from the model the accuracy drops to 0.98 and the F1 score drops to 0.96. Meaning that wind features do improve model performance.

Model	Accuracy	F1	AUC
No Wind	0.98	0.96	0.99
Just Wind	0.96	0.91	0.98

Model Accuracy F1 AUC

Obtaining the Data

For this part of the process I will use the dataset containing all homes.

```
In [2]: 1 #opening dataset
2 all_df = pd.read_csv(r'data\wind_modeling.csv')
3 all_df.describe()
```

Out[2]:

	AWND	WSF2	SizeRank	before	after	percent	i
count	344.000000	344.000000	344.000000	344.000000	344.000000	344.000000	344
mean	14.046948	27.354942	1359.034884	196091.514509	219108.234733	13.783402	(
std	7.013264	12.690591	3021.571059	118770.632169	126637.076629	8.838321	(
min	2.910000	0.000000	12.000000	44457.567490	52157.021900	-4.161603	(
25%	8.720000	18.100000	106.000000	101212.031625	122666.923050	6.526272	(
50%	12.750000	23.900000	233.000000	168065.082200	188091.092300	11.638427	(
75%	17.220000	31.325000	916.000000	257806.739650	286217.803975	19.627916	(
max	40.710000	79.000000	12877.000000	638691.713200	669502.004200	38.720009	1

Running the Model Without Wind Features

```
In [3]: 1 #initiating model
2 xgb = XGBClassifier(random_state=56)
```

Selecting Our Target Variable and Features

```
In [4]: 1 #y is prediction variable
2 #X is features
3 y_boost = all_df['increase']
4 X_boost = all_df.drop(['City', 'HurricaneName', 'DATE', 'after', 'perc
5
```

Train/Test Split

```
In [5]: 1 #train/test splits with 30% test size
2 XG_X_train, XG_X_test, XG_y_train, XG_y_test = train_test_split(X_boos
```

Training Data

```
In [6]: 1 #fitting the model
        2 xgb.fit(XG_X_train, XG_y_train);
```

```
In [7]: 1 #getting predictions
        2 y_pred_train = xgb.predict(XG_X_train)
```

```
In [8]: 1 #Printing Accuracy
        2 print('Accuracy: %.3f' % accuracy_score(XG_y_train, y_pred_train))
```

Accuracy: 0.996

```
In [9]: 1 #using F-1 score to see how it performs
        2 #F1 Score = 2* Precision Score * Recall Score/ (Precision Score + Recall Score)
        3 print('F1 Score: %.3f' % f1_score(XG_y_train, y_pred_train))
```

F1 Score: 0.992

Testing Data

```
In [10]: 1 #fitting the model
         2 xgb.fit(XG_X_test, XG_y_test);
```

```
In [11]: 1 #getting predictions
         2 y_pred_test = xgb.predict(XG_X_test)
```

```
In [12]: 1 #Printing Accuracy
         2 accuracy_XG_nowind = accuracy_score(XG_y_test, y_pred_test)
         3 print(accuracy_XG_nowind)
```

0.9807692307692307

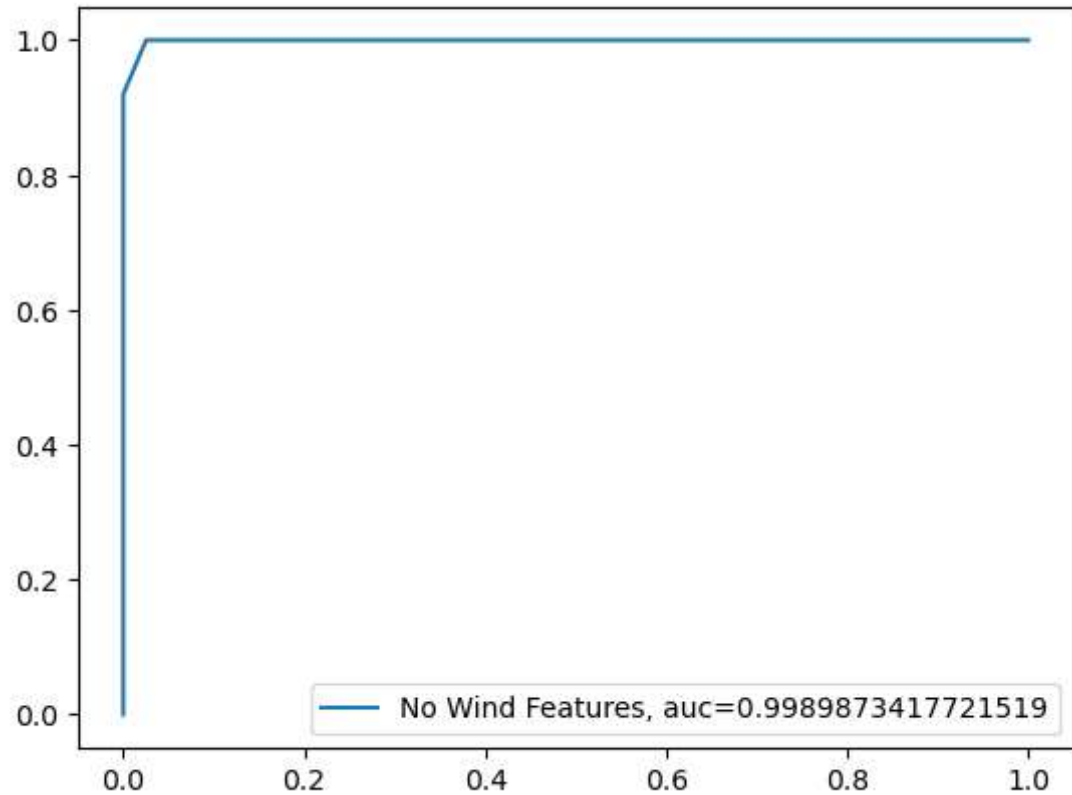
```
In [13]: 1 #using F-1 score to see how it performs
         2 #F1 Score = 2* Precision Score * Recall Score/ (Precision Score + Recall Score)
         3 F1_XG_nowind = f1_score(XG_y_test, y_pred_test)
         4 print(F1_XG_nowind)
```

0.96

```
In [14]: 1 model_dict = {}
         2 model_dict['XGBoost Accuracy No Wind'] = accuracy_XG_nowind
         3 model_dict['XGBoost F1 No Wind'] = F1_XG_nowind
         4 model_dict
```


```
Out[14]: {'XGBoost Accuracy No Wind': 0.9807692307692307, 'XGBoost F1 No Wind': 0.96}
```

```
In [15]: 1 #Let's check out the AUC curve
2 #getting probability
3 y_pred_prob = xgb.predict_proba(XG_X_test)[:,:1]
4 fpr, tpr, _ = metrics.roc_curve(XG_y_test, y_pred_prob)
5 auc = metrics.roc_auc_score(XG_y_test, y_pred_prob)
6
7 #plotting
8 plt.plot(fpr,tpr,label="No Wind Features, auc="+str(auc))
9 plt.legend(loc=4)
10 plt.show()
```

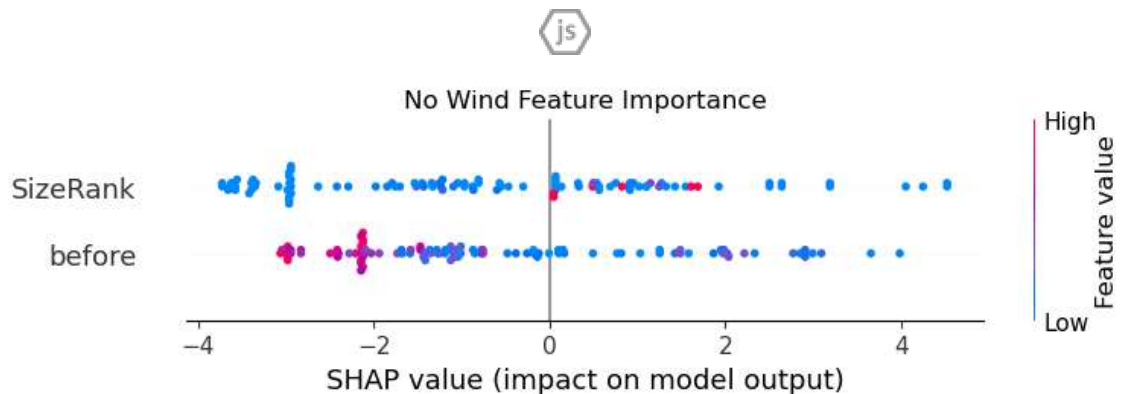


```
In [16]: 1 #plotting confusion matrix
2 conf_matrix = confusion_matrix(y_true=XG_y_test, y_pred=y_pred_test)
3
4 fig, ax = plt.subplots(figsize=(5, 5))
5 ax.matshow(conf_matrix, cmap=plt.cm.Blues, alpha=0.3)
6 for i in range(conf_matrix.shape[0]):
7     for j in range(conf_matrix.shape[1]):
8         ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center')
9
10 plt.xlabel('Predictions', fontsize=18)
11 plt.ylabel('Actuals', fontsize=18)
12 plt.title('Confusion Matrix for No Wind Features', fontsize=18)
13 plt.show()
```

Confusion Matrix for No Wind Features



```
In [17]: 1 #Using SHAP to assess feature importance
2 # creating an explainer for our model
3 explainer = shap.TreeExplainer(xgb)
4
5 # finding out the shap values using the explainer
6 shap_values = explainer.shap_values(XG_X_test)
7
8 #creating a beeswarm plot
9 shap.initjs()
10 plt.title("No Wind Feature Importance", y=1)
11 shap.summary_plot(shap_values, XG_X_test)
```



XGBoost on Just Wind Features

Let's see our results using just wind as a feature.

Selecting Our Target Variable and Features

```
In [18]: 1 #y is prediction variable
2 #X is features
3 y_boost = all_df['increase']
4 X_boost = all_df.drop(['City', 'HurricaneName', 'DATE', 'after', 'perc
5
```

Train/Test Split

```
In [19]: 1 #train/test splits with 30% test size
2 XG_X_train, XG_X_test, XG_y_train, XG_y_test = train_test_split(X_boos
```

Training Data

```
In [20]: 1 #fitting model
2 xgb.fit(XG_X_train, XG_y_train);
```

```
In [21]: 1 #getting predictions
        2 y_pred_train = xgb.predict(XG_X_train)
```

```
In [22]: 1 #Printing Accuracy
        2 print('Accuracy: %.3f' % accuracy_score(XG_y_train, y_pred_train))

Accuracy: 0.933
```

```
In [23]: 1 #using F-1 score to see how it performs
        2 #F1 Score = 2* Precision Score * Recall Score/ (Precision Score + Recall Score)
        3 print('F1 Score: %.3f' % f1_score(XG_y_train, y_pred_train))

F1 Score: 0.857
```

Testing Data

```
In [24]: 1 #fitting the model
        2 xgb.fit(XG_X_test, XG_y_test);
```

```
In [25]: 1 #getting predictions
        2 y_pred_test = xgb.predict(XG_X_test)
```

```
In [26]: 1 #Printing Accuracy
        2 accuracy_XG_justwind = accuracy_score(XG_y_test, y_pred_test)
        3 print(accuracy_XG_justwind)

0.9615384615384616
```

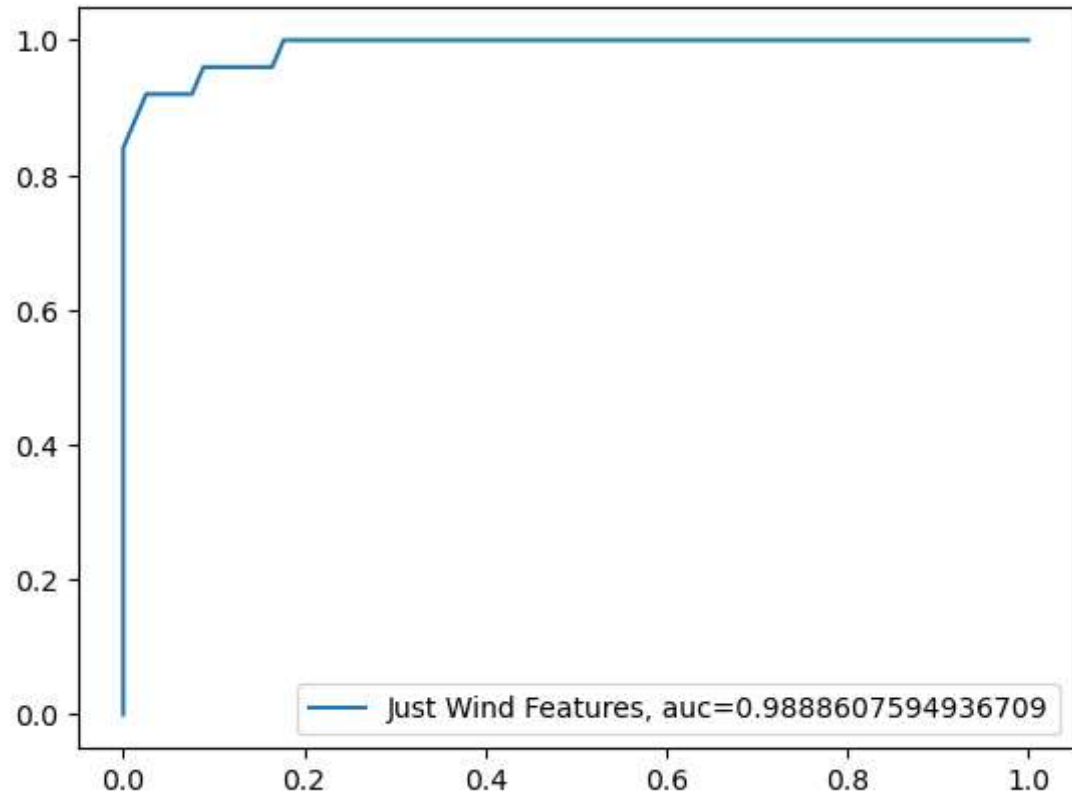
```
In [27]: 1 #using F-1 score to see how it performs
        2 #F1 Score = 2* Precision Score * Recall Score/ (Precision Score + Recall Score)
        3 F1_XG_justwind = f1_score(XG_y_test, y_pred_test)
        4 print(F1_XG_justwind)

0.9166666666666666
```

```
In [28]: 1 model_dict['XGBoost Accuracy Just Wind'] = accuracy_XG_justwind
        2 model_dict['XGBoost F1 Just Wind'] = F1_XG_justwind
        3 model_dict
```


```
Out[28]: {'XGBoost Accuracy No Wind': 0.9807692307692307,
          'XGBoost F1 No Wind': 0.96,
          'XGBoost Accuracy Just Wind': 0.9615384615384616,
          'XGBoost F1 Just Wind': 0.9166666666666666}
```

```
In [29]: 1 #Let's check out the AUC curve
2 #getting probability
3 y_pred_prob = xgb.predict_proba(XG_X_test)[::,1]
4 fpr, tpr, _ = metrics.roc_curve(XG_y_test, y_pred_prob)
5 auc = metrics.roc_auc_score(XG_y_test, y_pred_prob)
6
7 #plotting
8 plt.plot(fpr,tpr,label="Just Wind Features, auc="+str(auc))
9 plt.legend(loc=4)
10 plt.show()
```

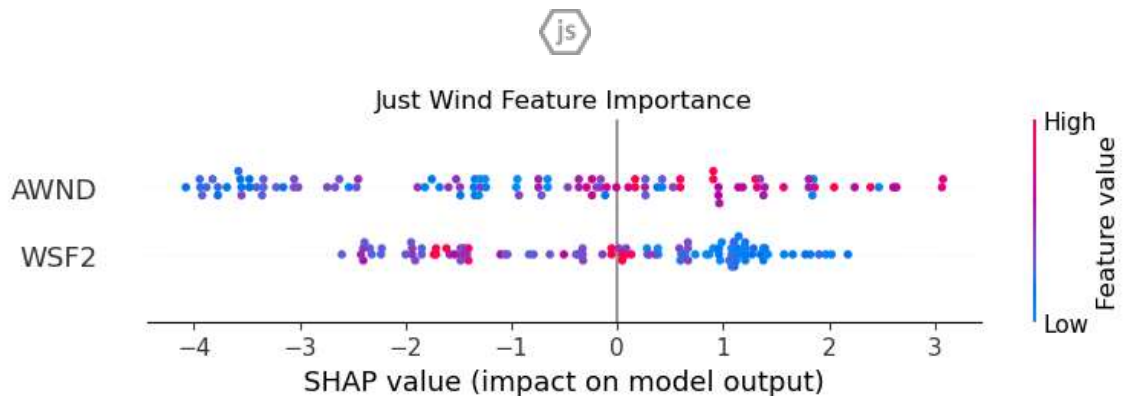



```
In [30]: 1 #plotting confusion matrix
2 conf_matrix = confusion_matrix(y_true=XG_y_test, y_pred=y_pred_test)
3
4 fig, ax = plt.subplots(figsize=(5, 5))
5 ax.matshow(conf_matrix, cmap=plt.cm.Blues, alpha=0.3)
6 for i in range(conf_matrix.shape[0]):
7     for j in range(conf_matrix.shape[1]):
8         ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center')
9
10 plt.xlabel('Predictions', fontsize=18)
11 plt.ylabel('Actuals', fontsize=18)
12 plt.title('Confusion Matrix of Just Wind Features', fontsize=18)
13 plt.show()
```

Confusion Matrix of Just Wind Features



```
In [31]: 1 #Using SHAP to assess feature importance
2 # creating an explainer for our model
3 explainer = shap.TreeExplainer(xgb)
4
5 # finding out the shap values using the explainer
6 shap_values = explainer.shap_values(XG_X_test)
7
8 #creating a beeswarm plot
9 shap.initjs()
10 plt.title("Just Wind Feature Importance", y=1)
11 shap.summary_plot(shap_values, XG_X_test)
```



XG Boost Tuning

Let's see if we can improve our results using just wind features with model tuning. The optimal parameters according to the gridsearch are:

Best Parameters for Our XGBoost Model:

```
{'gamma': 0,
 'learning_rate': 0.01,
 'max_depth': 2,
 'min_child_weight': 1,
 'n_estimators': 10,
 'subsample': 0.5}
```

Selecting Our Target Variable and Features

```
In [32]: 1 #y is prediction variable
2 #X is features
3 y_boost = all_df['increase']
4 X_boost = all_df.drop(['City', 'HurricaneName', 'DATE', 'after', 'perc
5
```

Train/Test/Tune Split

In order to properly tune our model it should be done on data that is not seen during the training or testing process. This will help assure an unbiased model. 70% of our data will be used for training, 15% for testing, and 15% for tuning.

Works Cited:

1.Training, Tuning, and Test Datasets. onepager.togaware.com. Accessed July 6, 2023.

https://onepager.togaware.com/Training_Tuning_Test_Datase.html

(https://onepager.togaware.com/Training_Tuning_Test_Datase.html)

```
In [33]: 1 #train/test splits with 30% test size
2 #half of the test set will be used to create the set we will be tuning
3 XG_X_train, XG_X_test_tune, XG_y_train, XG_y_test_tune = train_test_sp
```

```
In [34]: 1 #train/test splits with 30% test size
2 XG_X_test, XG_X_tune, XG_y_test, XG_y_tune = train_test_split(XG_X_tes
```

```
1 #skipping this cell from running since we know the parameters
2
3 #hypertuning the model using GridSearch
4 #the parameters that will be searched
5 xgb_grid = {
6     'learning_rate': [0.01, 0.1, 0.5],
7     'gamma': [0, 0.01, 0.1],
8     'max_depth': [2, 5, 6, 10],
9     'min_child_weight': [0.1, 1, 10],
10    'subsample': [0.5, 0.7, 0.9],
11    'n_estimators': [5, 10, 20, 100]
12 }
13 #running GridSearch
14 xgb_gridsearch = GridSearchCV(estimator=xgb,
15                               param_grid=xgb_grid,
16                               cv=5,
17                               return_train_score=True)
18 #fitting the gridsearch
19 xgb_gridsearch.fit(XG_X_tune, XG_y_tune)
20
21 #printing the best parameters
22 print('Best Parameters for Our XGBoost Model:')
23 xgb_gridsearch.best_params_
24
```

```
In [35]: 1 #initiating the model
2 xgb_gs = XGBClassifier(gamma= 0, learning_rate= 0.01, max_depth= 2, mi
```

Training Data

```
In [36]: 1 #fitting the model
2 xgb_gs.fit(XG_X_train, XG_y_train);
```

```
In [37]: 1 #predicting
2 y_pred_train = xgb_gs.predict(XG_X_train)
```

```
In [38]: 1 #Printing Accuracy
2 print('Accuracy: %.3f' % accuracy_score(XG_y_train, y_pred_train))
```

Accuracy: 0.746

```
In [39]: 1 #using F-1 score to see how it performs
2 #F1 Score = 2* Precision Score * Recall Score / (Precision Score + Recall Score)
3 print('F1 Score: %.3f' % f1_score(XG_y_train, y_pred_train))
```

F1 Score: 0.000

Testing Data

```
In [40]: 1 #fitting the model
2 xgb_gs.fit(XG_X_test, XG_y_test);
```

```
In [41]: 1 #getting predictions
2 y_pred_test = xgb_gs.predict(XG_X_test)
```

```
In [42]: 1 #Printing Accuracy
2 accuracy_XG_justwind_tuned = accuracy_score(XG_y_test, y_pred_test)
3 print(accuracy_XG_justwind_tuned)
```

0.7884615384615384


```
In [43]: 1 #using F-1 score to see how it performs
2 #F1 Score = 2* Precision Score * Recall Score / (Precision Score + Recall Score)
3 F1_XG_justwind_tuned = f1_score(XG_y_test, y_pred_test)
4 print(F1_XG_justwind_tuned)
```

0.56

```
In [44]: 1 #adding values to model dictionary
2 model_dict['XGBoost Accuracy Just Wind Tuned'] = accuracy_XG_justwind_tuned
3 model_dict['XGBoost F1 Just Wind Tuned'] = F1_XG_justwind_tuned
4 model_dict
```

```
Out[44]: {'XGBoost Accuracy No Wind': 0.9807692307692307,
'XGBoost F1 No Wind': 0.96,
'XGBoost Accuracy Just Wind': 0.9615384615384616,
'XGBoost F1 Just Wind': 0.9166666666666666,
'XGBoost Accuracy Just Wind Tuned': 0.7884615384615384,
'XGBoost F1 Just Wind Tuned': 0.56}
```

```
In [45]: 1 #Let's check out the AUC curve
2 #getting probability
3 y_pred_prob = xgb_gs.predict_proba(XG_X_test)[::,1]
4 fpr, tpr, _ = metrics.roc_curve(XG_y_test, y_pred_prob)
5 auc = metrics.roc_auc_score(XG_y_test, y_pred_prob)
6
7 #plotting
8 plt.plot(fpr,tpr,label="Tuned Model, auc="+str(auc))
9 plt.legend(loc=4)
10 plt.show()
```



```
In [46]: 1 #plotting confusion matrix
2 conf_matrix = confusion_matrix(y_true=XG_y_test, y_pred=y_pred_test)
3
4 fig, ax = plt.subplots(figsize=(5, 5))
5 ax.matshow(conf_matrix, cmap=plt.cm.Blues, alpha=0.3)
6 for i in range(conf_matrix.shape[0]):
7     for j in range(conf_matrix.shape[1]):
8         ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center')
9
10 plt.xlabel('Predictions', fontsize=18)
11 plt.ylabel('Actuals', fontsize=18)
12 plt.title('Confusion Matrix for Tuned Model', fontsize=18)
13 plt.show()
```

